

# Introduction to Kubernetes

---

## Introduction to Kubernetes

### 1. Evolution of Software Architecture

#### 1.1 Monolithic Architecture

Pros

Cons

#### 1.2 Service-oriented Architecture (SOA)

Pros:

Cons:

#### 1.3 Microservice Architecture

Pros

Cons

### 2. Containerization and docker

Virtualization vs containers

Containerization

Linux Namespace

Linux Cgroups

UnionFS

### 3. Need for container orchestration tool

Challenges from manage multi containers

### 4. What is Kubernetes

Features provided by Kubernetes

### 5. Kubernetes Architecture

Components of the control plane

kube-apiserver

kube-scheduler

kube-controller-manager

etcd

Components of the worker nodes

Nodes

Pods

kubelet

kube-proxy

Container runtime

Addons

### 6. Basic commands in `kubectl`

The workflow when a change is requested

### 7. Kubernetes Objects

Imperative vs Declarative

Kubernetes Objects

How to Describe a Kubernetes object

TypeMeta

Metadata

Spec

Status

Demo: An example of object definition file

An overview of Kubernetes objects by groups

## 8. Deploying applications using Kubernetes

Node

Pod

Service

Ingress

ConfigMap & Secret

Volume

Deployment

## 9. Demo

Create Config Maps Imperative

Create Config Maps Imperative

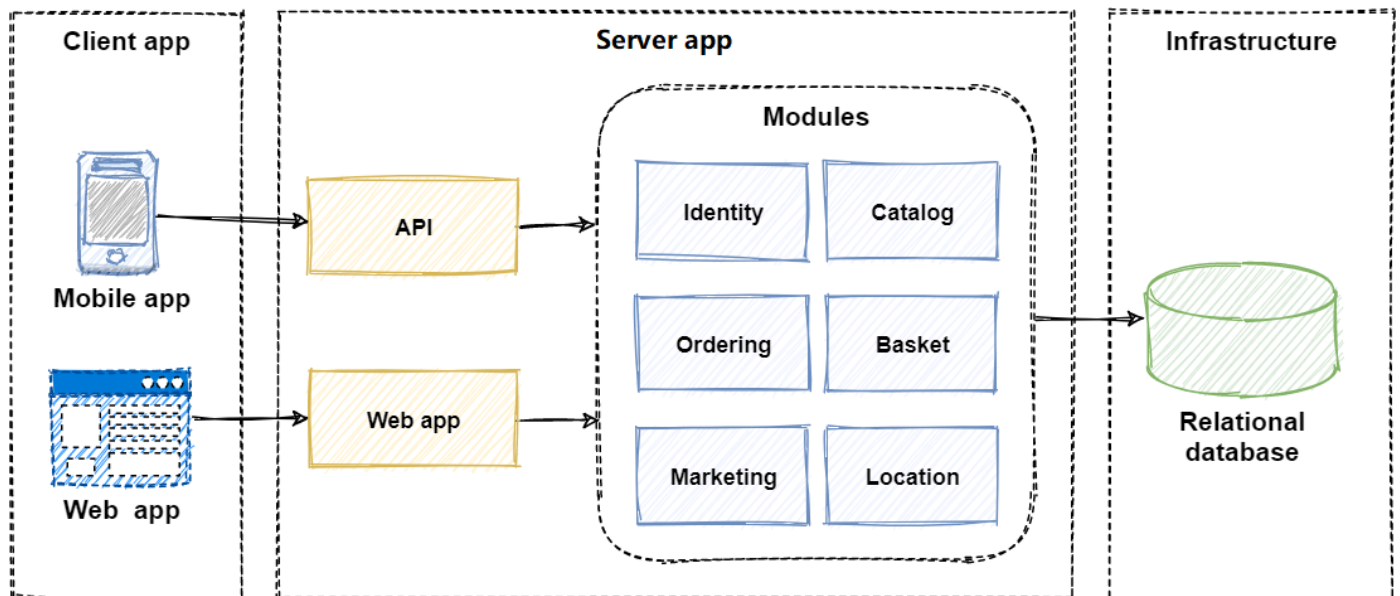
Use ConfigMaps in Pods

Deploying Wordpress and MySQL

References

# 1. Evolution of Software Architecture

## 1.1 Monolithic Architecture



The concept of monolithic model lies in different components of an application being combined into a single program on a single platform.

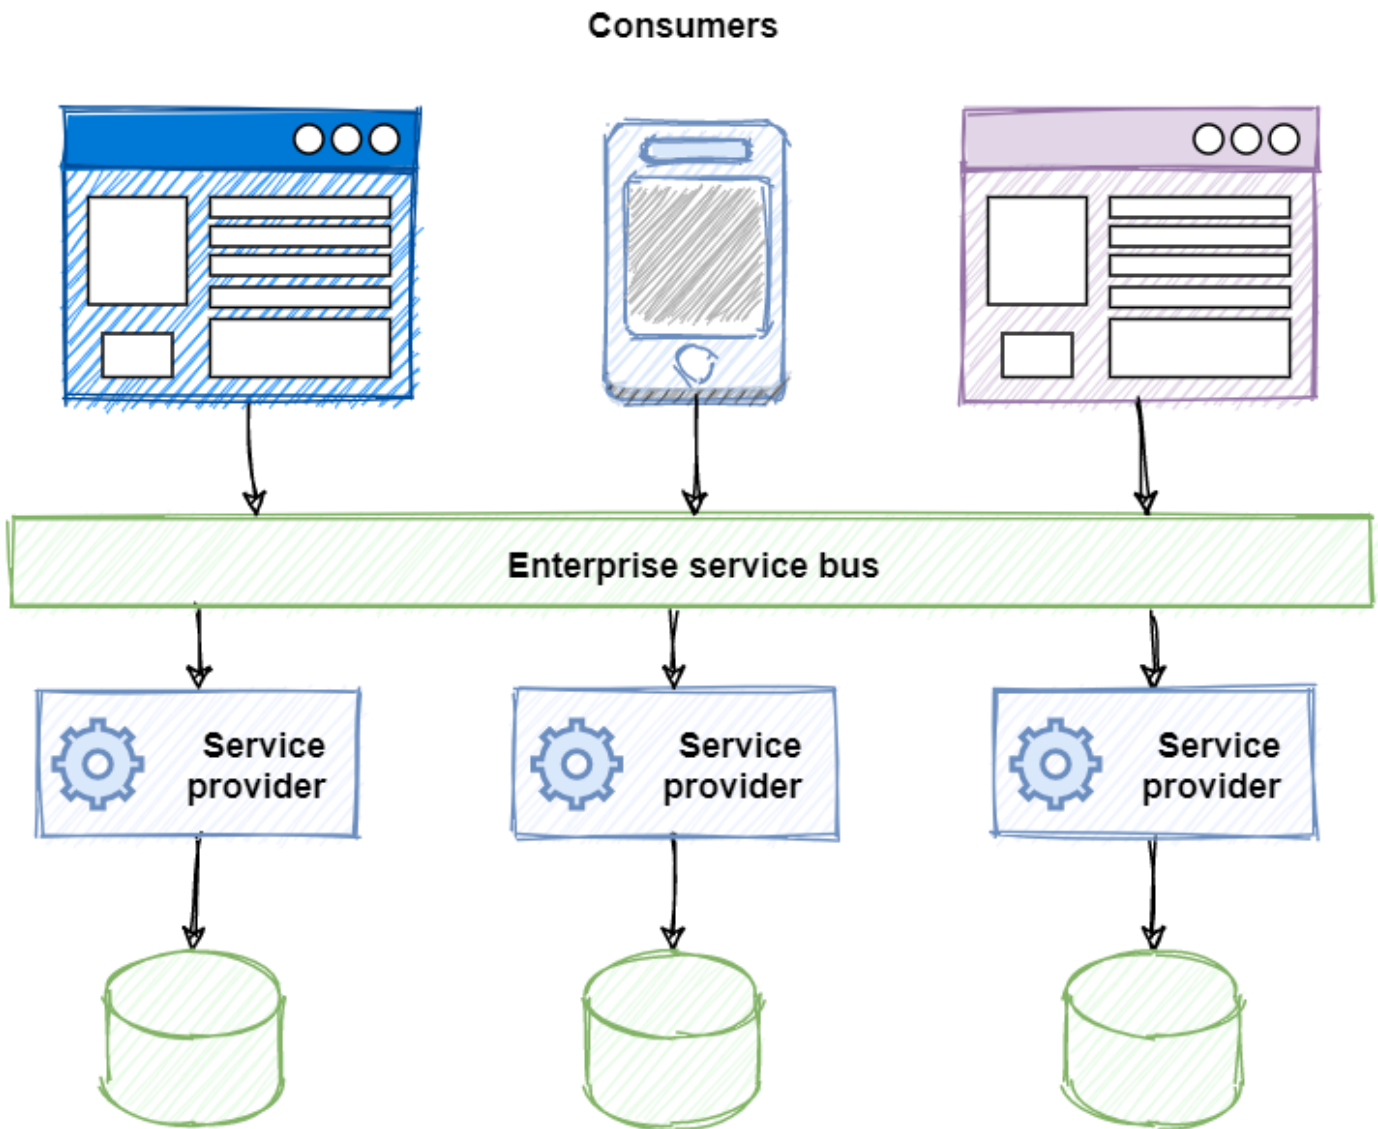
## **Pros**

- Simpler development and deployment
- Fewer cross-cutting concerns
- Better performance
- Perfect for small projects.

## **Cons**

- Codebase gets cumbersome overtime
- The monolith is becoming so complicated that no single person can understand it.
- Difficult to adopt new technologies
- Limited agility
- One unstable component can crash the entire system.

## **1.2 Service-oriented Architecture (SOA)**



In this model, **service consumers** and **service providers** collaborated via middleware messaging components, often referred to as an **Enterprise Service Bus**.

- Centralized service providers registered with the ESB
- Business logic would be built into the ESB to integrate providers and consumers.
- Service consumers could then find and communicate with these providers using the ESB.

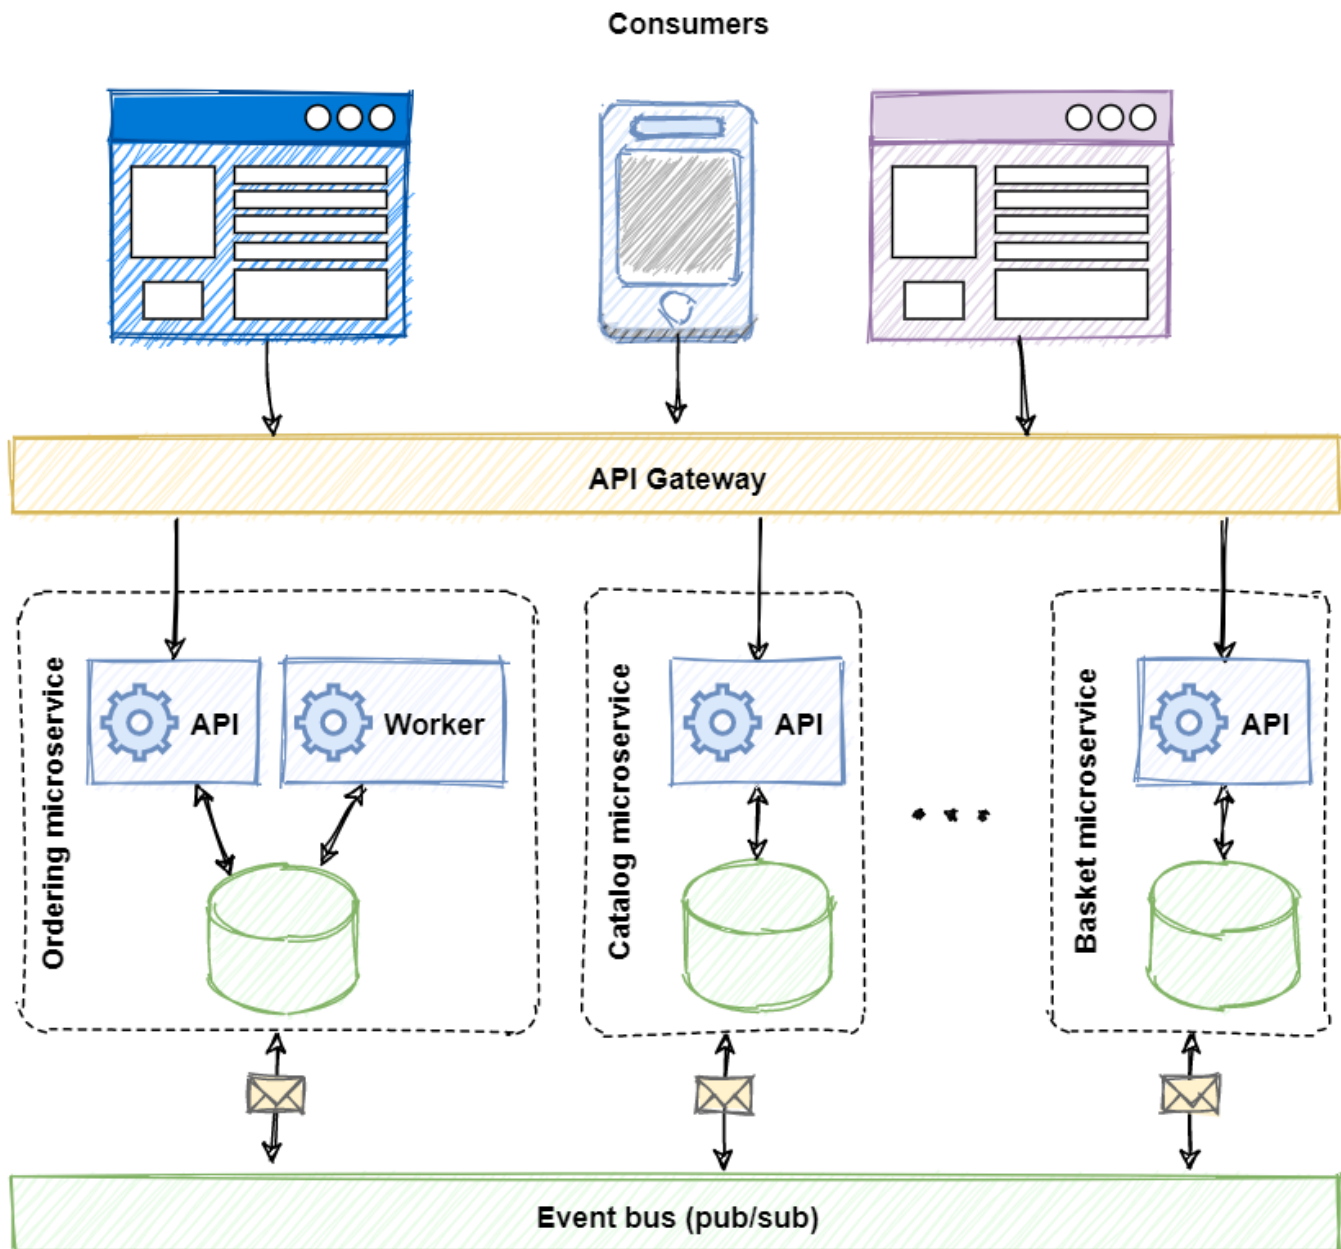
#### Pros:

- Reusability of services
- Better maintainability
- Higher reliability
- Parallel development

## Cons:

- Complex management
- Maintenance costs became high and ESB middleware expensive.
- Services tended to be large. They often shared dependencies and data storage.
- Extra overload: when using multiple services, this increases response time and decreases overall performance.

## 1.3 Microservice Architecture



Microservice apps consist of **multiple independent components** that are glued together with **APIs**.

The microservices approach focuses mainly on **business priorities and capabilities**.

## Pros

- Easy to develop, test, and deploy
- Increased agility: several teams can work on their services independently and quickly
- Ability to scale horizontally

## Cons

- Increased demand for automation, as every service should be tested and monitored
- Data consistency and transaction management becomes harder as each service has a database
- Security concerns
- Different programming languages

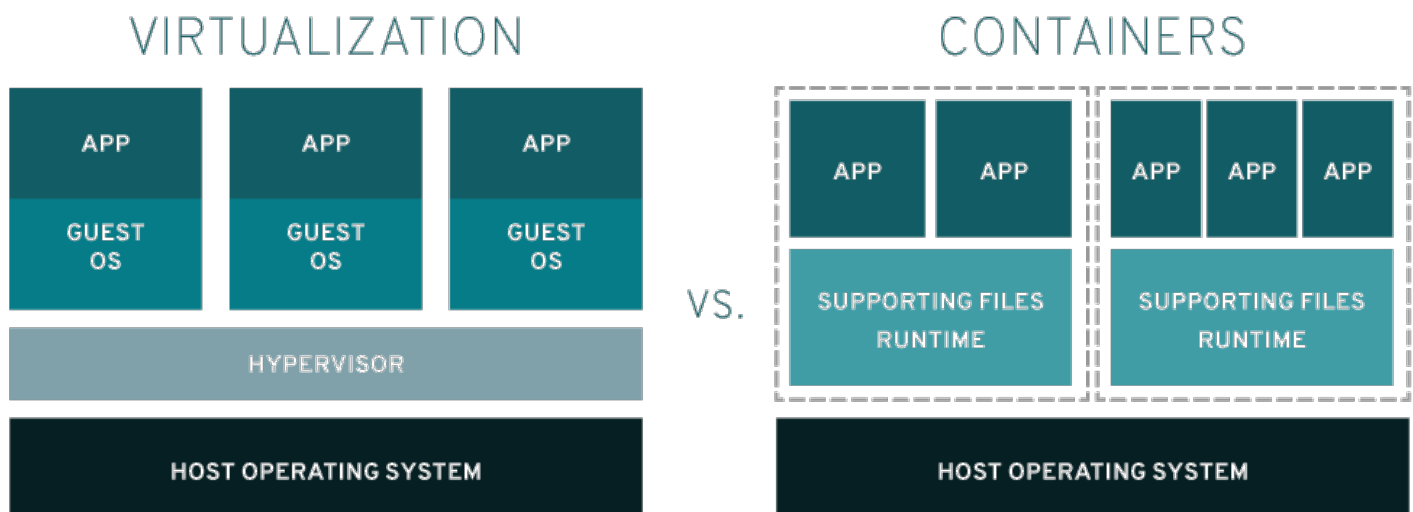
## 2. Containerization and docker

---

Microservices architecture presents a set of challenges

- Multiple Network Calls among services
- Multiple microservices deployed on the same host, services are less isolated

## Virtualization vs containers



## Containerization

**Containerization** packages together everything needed to run a single application or microservice (along with runtime libraries they need to run). The container **includes all the code, its dependencies** and even the operating system itself.

High level container runtimes , such as **Docker** simplify the process of creating and maintaining containers.

## Linux Namespace

- A feature provided by Linux Kernel to isolate resources.
- For containers, a namespace is what **defines the boundaries of a process'** "awareness" of what else is running around it.
- Namespace resources: IPC, Network, PID, Mount, PTS, USR.

## Linux Cgroups

- The mechanism provided by Linux Kernel to control or monitor one or more processes.
- Control CPU, Memory, I/O of processes.
- Cgroup driver: systemd as init system will be the cgroup manager, but Docker uses cgroupfs as driver by default.

## UnionFS

- Unite several directories into a single virtual filesystem
- Set readonly, readwrite, whiteout-able permission for each directory
- Layered File System.

## 3. Need for container orchestration tool



## Challenges from manage multi containers

- Increased usage of containers. In production, you can end up with dozens, even thousands of containers over time.
- Demand for a proper way managing those containers
- Integrate and orchestrate these modular parts
- Scale up and scale down based on the demand
- Make them fault tolerant



- Provide communication across a cluster

**Container orchestration** is the automated process of coordinating and organizing all aspects of individual containers, their functions, and their dynamic environments.

Using container orchestration, one can

- manage when and how containers start and stop,
- schedule and coordinate components' activities,
- monitor health,
- institute failover and recovery processes.

## 4. What is Kubernetes

Kubernetes aka. k8s (pronounced /kates/)



Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation.

- Open source container orchestration tool
- Originally developed by Google
- First graduated CNCF project
- Open source version of Borg
- Helps manages containerized applications in different environments

## Features provided by Kubernetes

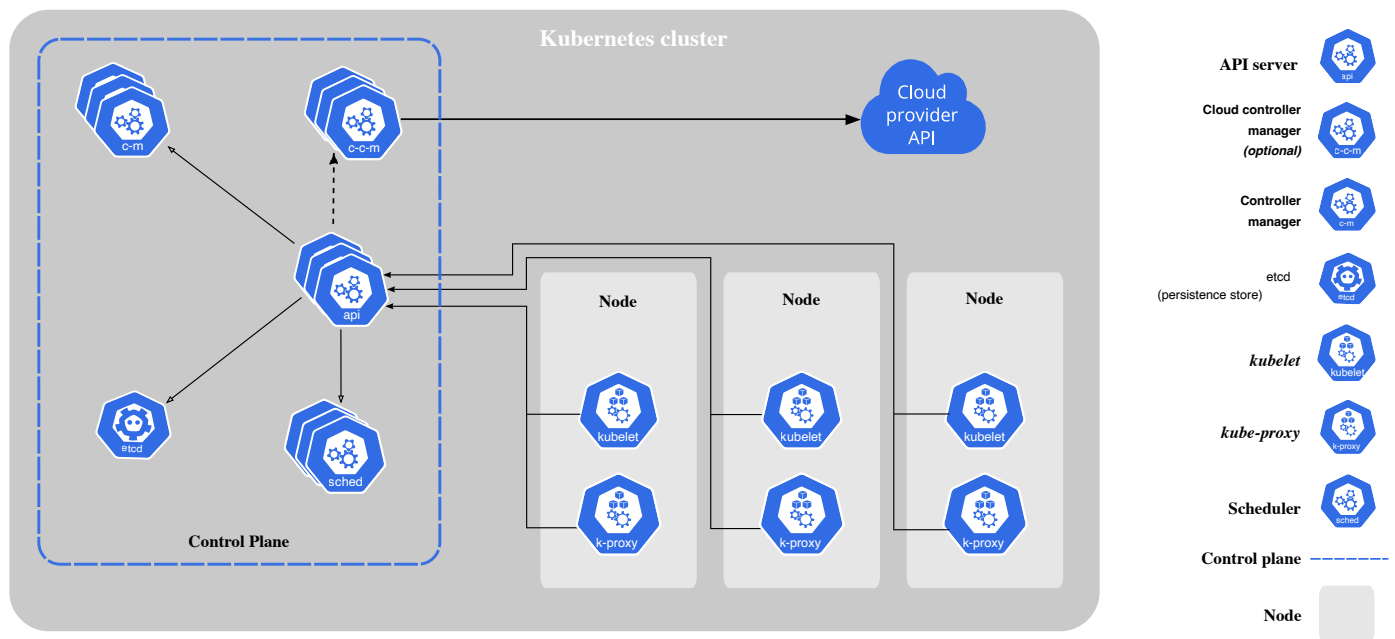
- **High Availability**
  - Application High Availability building on ReplicaSet, StatefulSet
  - Kubernetes Components are high available
- **Security**
  - Services are running on TLS
  - Provide ServiceAccount and User Account
  - Isolation by Namespace



- Secret
- Taints, PSP, Network Policy
- **Portability**
  - multi host os
  - Multi cloud and hybrid cloud
- **Extensibility**
  - CRD
  - Plugins (Addons)

## 5. Kubernetes Architecture

You can visualize a Kubernetes cluster as two parts: the **control plane** and **the worker machines, or nodes**.



### Components of the control plane

Control plane is the nerve center of our Kubernetes cluster.

- Control the cluster, along with data about the cluster's state and configuration.
- Handle the important work of making sure your containers are running in sufficient numbers and with the necessary resources.

#### kube-apiserver

- The Kubernetes API is the front end of the Kubernetes control plane, **handling internal and external requests**.
- You can access the API through **REST** calls, through the **kubectl** command-line interface, or through other command-line tools such as kubectl.

#### kube-scheduler

## kube-scheduler

- Control plane component that **watches for newly created Pods** with no assigned node, and selects a node for them to run on.
- The scheduler considers the **resource needs of a pod**, such as CPU or memory, along with the **health** of the cluster.
- It is only responsible for deciding which pod goes on which node. **It doesn't actually place the pod on the nodes.** The **kubelet** or the captain on the ship is who creates the pod on the ships.

## kube-controller-manager

- Controllers take care of actually running the cluster
- Kubernetes controller-manager contains several controller functions in one.
  - **Node-Controller**
  - **Replication-Controller:** It is responsible for monitoring the status of replicaset and ensuring that the desired number of PODs are available at all times within the set.
  - PV-Binder-Controller
  - Service-Account-Controller
  - StatefulSet
  - ReplicaSet
  - CronJob
  - Job-Controller
  - PV-Protection-Controller
  - Deployment-Controller
  - Namespace-Controller
  - Endpoint-Controller
- All of these controllers are packaged into a single process called `Kube-Controller-Manager`.

## etcd

- A key value database for Cluster Data Store
- Configuration data and information about the state of the cluster.

## Components of the worker nodes

Node components run on every node, maintaining running pods and providing the Kubernetes runtime environment.

### Nodes

It is an abstraction for the compute node. Every cluster has at least one worker node to run containerized applications.

### Pods

## Pods

The ultimate aim is to **deploy** our application in the form of **containers** on a set of machines that are **configured as worker nodes** in a cluster.

However, **Kubernetes does not deploy containers directly on the worker nodes.**

- A pod is **the smallest and simplest unit** in the Kubernetes object model.
- It represents **a single instance of an application.**
- Each pod is made up of a container or a series of tightly coupled containers

## kubelet

- An agent that runs on each node in the cluster. It makes sure that containers are running in a Pod
- It communicates with the control plane
- It will not maintain containers not created by Kubernetes
- Jobs performed by kubelet: Register Node, Create Pods, Monitor Node and Pod

## kube-proxy

- A **Pod network** is a virtual network that **spans across all the nodes in the cluster** to which all the Pods connect to.
- Each compute node also contains kube-proxy, a network proxy for facilitating Kubernetes networking services.
- kube-proxy maintains network rules on nodes.
- Works on the Level 4 TCP/IP stack
- Underhood is iptables or ipvs
  - Creates an IP table rule on each node in the cluster.
  - Forwarding traffic heading to the IP of the service to the IP of the actual Pod.
- It's job is to **look for new services** and every time a new service is created it creates the appropriate rules on each node to **forward traffic to those services to the backend pods.**

## Container runtime

To run the containers, each compute node has a container runtime engine. For example, Docker.

## Addons

- kube-dns: provide DNS Service for the cluster
- Ingress Controller: provide external access for service
- Metrics Server: provide resource monitoring
- Dashboard: provide web-based UI

```
# View component definition files in kubeadm
ls /etc/kubernetes/manifests
```

## 6. Basic commands in `kubectl`

---

When running a `kubectl` command,

- The `kubectl` utility is in fact reaching to the kube-api server.
- The kube-api server first authenticates the request and validates it.
- It then retrieves the information from the etcd cluster and response back to the user.

### The workflow when a change is requested

```
curl -X POST /api/v1/namespaces/default/pods ...
```

1. Authenticate user
2. Validate request
3. Retrieve data
4. Update ETCD
5. Scheduler
6. Kubelet

- The scheduler continuously monitors the API server and realizes that there is a new pod with no node assigned.
- The scheduler identifies the right node to place the new POD on and communicates that back to the kube-api server.
- The API server then updates the information in the ETCD cluster.
- The API server then passes that information to the kubelet in appropriate worker node.
- The kubelet then creates the POD on the node and instructs the container runtime engine to deploy the application image.
- Once done, the kubelet updates the status back to the API server and the API server then updates the data back in the ETCD cluster.

```
kubectl get po -oyaml -w
```

```
kubectl describe po xxx
```

```
kubectl exec
```

```
kubectl logs
```

## 7. Kubernetes Objects

---

### Imperative vs Declarative

# Imperative vs Declarative



Imperative vs Declarative



- Everything in Kubernetes is abstracted as an object.
- Idempotency
- API Standardization

## Kubernetes Objects

- Kubernetes objects are **persistent entities** in the Kubernetes system.
  - Kubernetes uses these entities to **represent the state of your cluster**.
  - By creating an object, you are specifying the cluster's **desired state**.
  - Kubernetes provides **Kubernetes API** to manage these objects.
- The API object is the **base unit** to manage a Kubernetes Cluster.
- New API objects will be introduced when new features are added to Kubernetes Cluster

```
kubectl api-resources
```

```
kubectl explain
```

## How to Describe a Kubernetes object

Every kubernetes object has below four types of properties

### TypeMeta

The most fundamental definition of an object, uses **GKV(Group, Kind, Version)** to define the type of an object.

- **Group:** The kubernetes API is grouped into multiple such groups based on their purposes, `/metrics`, `/healthz`, `/version`, `/api`, `/apis`, `/logs`.
  - Core: `/api`: namespaces, pods, replication controllers, events, endpoints, nodes, bindings, persistent volumes, persistent volume claims, configmaps, secrets, services, etc.
  - Named: `/apis`: apps, extensions, networking, storage, authentication certificates etc.
- **Kind:** Defines the type of the object, for example, Node, Pod, Deployment, Service, etc.
- **Version:** Different API versions indicate different levels of stability and support.

### Metadata

## Metadata

- **Namespace** and **Name** are the most important two properties of Metadata type. They uniquely identifies an instance of an object.
- **Label** defines whether the object can be recognized. An object can have any number of labels in key-value format. Supports filters.
- **Annotation** is also in key-value format, but serve different purposes: for example, assist deploying application, security policy and scheduling policy. Rolling update information in deployment.
- **Finalizer**: a resource lock. When deleting an object and the resource lock is not null, then it's an logical delete.
- **ResourceVersion**: Every object has an resource object at any time.

## Spec

- providing a description of the characteristics you want the resource to have, i.e. the **desired state**.
- different objects may have different specs.
- Defined by the client

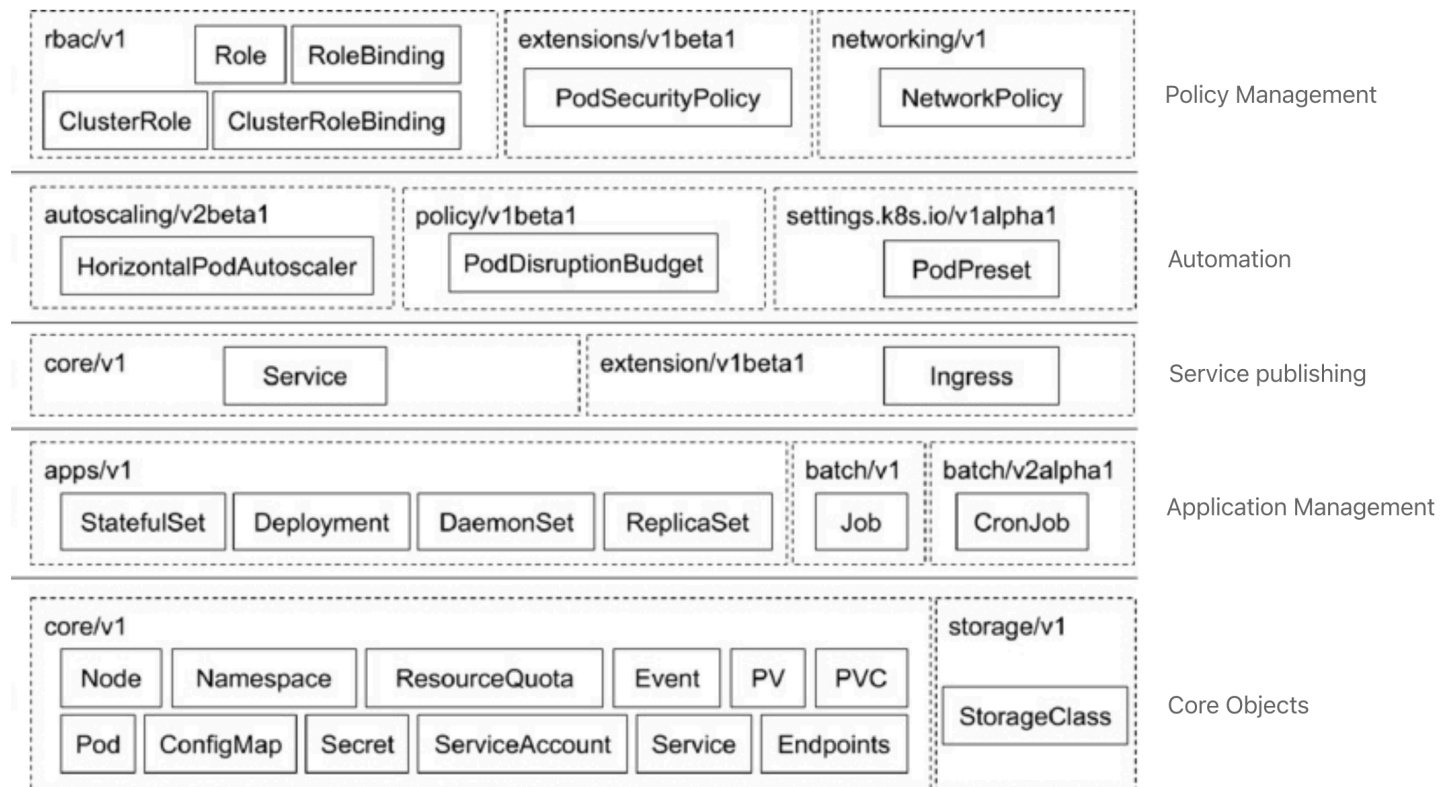
## Status

- describes the **current state** of the object
- It's updated by the controllers.
- The **control plane** continually and actively **manages every object's actual state** to match the desired state

## Demo: An exmaple of object definition file

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```


## An overview of Kubernetes objects by groups



## 8. Deploying applications using Kubernetes



# Cluster



## Node

- host machine, can be physical machine or virtual machine

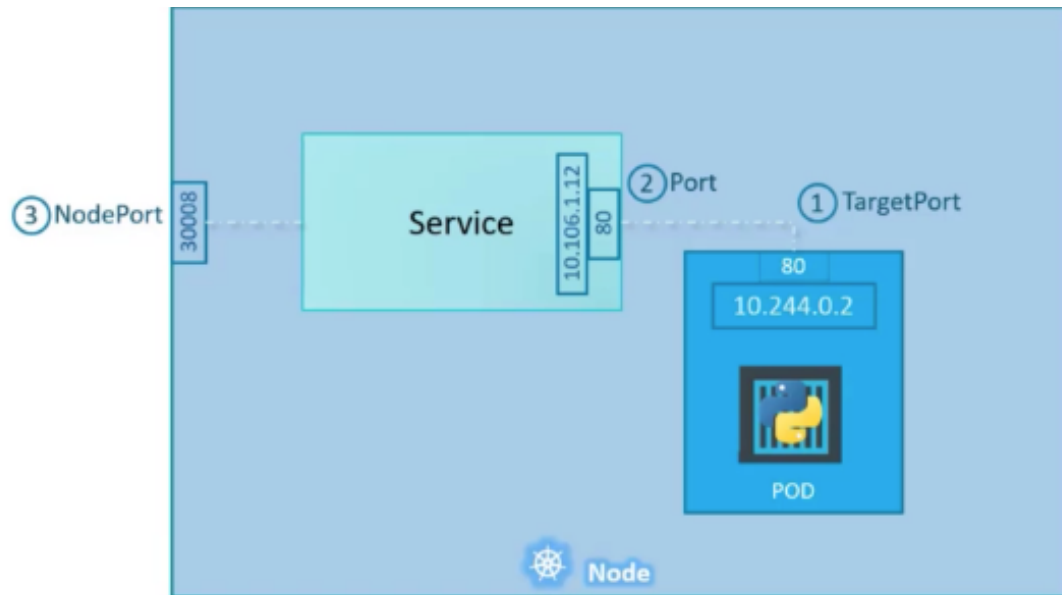
## Pod

- Smallest unit in kubernetes for scheduling
- An abstraction over container
- Contains a group of containers, normally one application per pod
- Sharing PID, IPC, Network, File System etc.
- Each Pod gets its own IP address
- Pods are ephemeral, each time new Pod are created, they are assigned a new IP.

## Service

- Service is static or permanent IP address.
- Services enable communication between various components within and outside of the application.
- Every Pod can be attached to a service. Different lifecycle compared to Pod.
- By default, The service is in fact like **a virtual server inside the node**. the service has its own IP address and that IP address is called **the cluster IP of the service**. (DB Service)
- Expose **external service** (Web App Service)
  - NodePort: **makes an internal Pod accessible** on a Port on the Node.

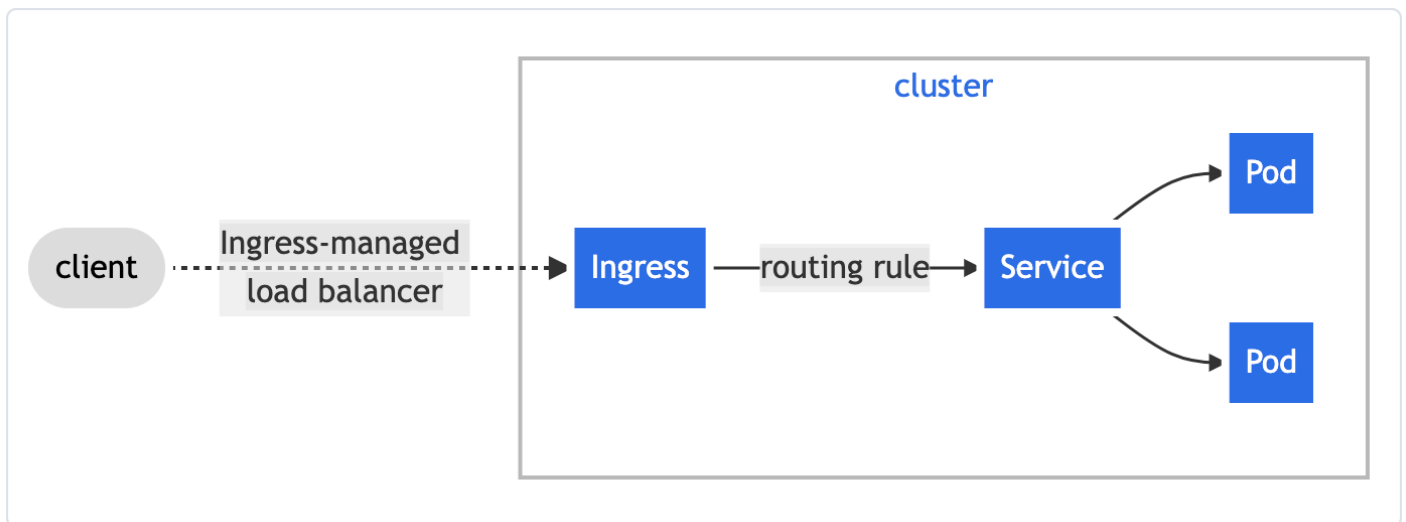
o



- Load Balancing

## Ingress

- manages external access to the services in a cluster
- exposes HTTP and HTTPS routes from outside the cluster to services within the cluster.



- Works on level 7 HTTP stack

## ConfigMap & Secret

- **ConfigMaps** are used to **pass configuration data in the form of key value pairs** in Kubernetes.e.g. service url
  - o Create ConfigMap
  - o Inject into Pod
- **Secrets** are used to **store sensitive information** like passwords or keys.

They are similar to ConfigMap except that they're **stored in an encoded or hashed format**.

## Volume

- Mount external volume to Pod
- Volume: define the source of volume
- VolumeMounts: define how the volume mount to Pod

## Deployment

The deployment provides us with the capability to **upgrade the underlying instances** seamlessly using rolling updates, undo changes, and pause and resume changes as required.

- Provides declarative updates to applications.
- Scale up and down by specifying replicas
- The Kubernetes deployment object lets you:
  - Deploy a replica set or pod
  - Update pods and replica sets
  - Rollback to previous deployment versions
  - Scale a deployment
  - Pause or continue a deployment

## 9. Demo

---

```
# 1. Pod
# Create nginx yaml
kubectl run webapp --image=nginx --dry-run=client -o yaml > web-pod.yaml

# 2. Service
# Create nginx svc yaml
kubectl expose pod webapp --name=webappsvc --port=80 --target-port=80 --dry-run=client -o
yaml > web-svc.yaml
kubectl delete svc webappsvc

# 3. Deploy
# Create Objects
kubectl create deployment --image=nginx webappdeploy --dry-run=client -o yaml > web-
deploy.yaml

# Update Objects
kubectl edit deployment webappdeploy
kubectl scale deployment webappdeploy --replicas=3
kubectl set image deployment nginx nginx=nginx:1.18
```

## Create Config Maps Imperative

## Create Config Maps Imperative

```
# cm short for configmap
kubectl create cm

# simple ones using --from-literal option
kubectl create configmap \
    <config-name> --from-literal=<key>=<value>

kubectl create configmap app-config --from-literal=APP_COLOR=blue --from-
literal=APP_MOD=prod

# complicate ones using --from-file option
kubectl create configmap
    <config-name> --from-file=<path-to-file>

kubectl create configmap app-config --from-file=app_config.properties
```

## Create Config Maps Imperative

Create a `config-map.yaml`

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  APP_COLOR: blue
  APP_MODE: prod
```

```
kubectl create -f config-map.yaml
```

## Use ConfigMaps in Pods

Add the `envFrom` option in `pod-definition.yaml`.

`envFrom` is a list

```
apiVersion: v1
```

```

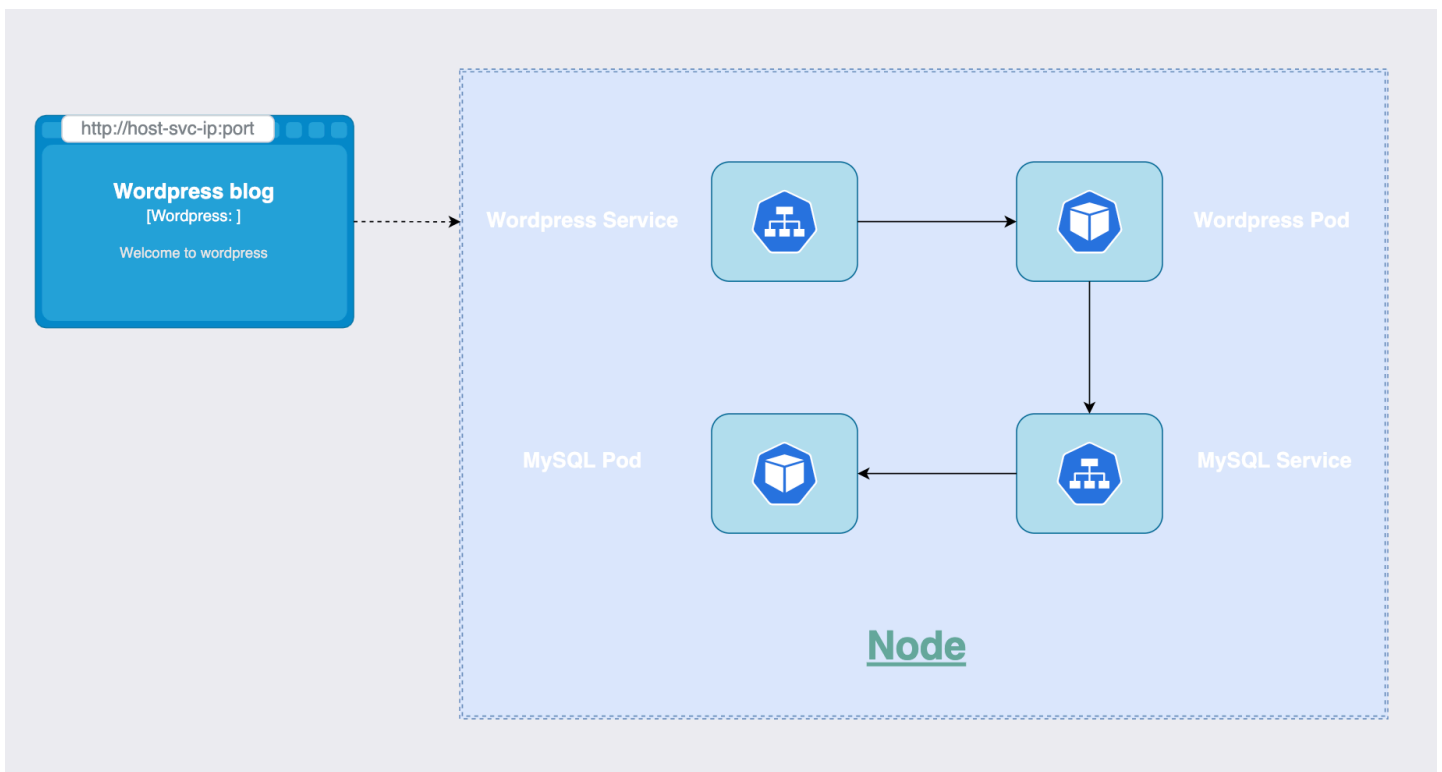
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp-color
spec:
  containers:
  - name: simple-webapp
    image: nginx
    ports:
    - containerPort: 8080
    envFrom:
    - configMapRef:
        name: app-config

```

To check the usage of envFrom:

```
kubectl explain pods --recursive | grep envFrom -A3
```

## Deploying Wordpress and MySQL



Steps to do this deployment

```

# Create mysql-secret for mysql Pod
kubectl create secret generic mysql-secret \
--from-literal=MYSQL_DATABASE=blog \
--from-literal=MYSQL_USER=admin \

```

```
--from-literal=MYSQL_PASSWORD=admin \
--from-literal=MYSQL_ROOT_PASSWORD=admin

# Create wordpress-secret for wordpress pod
kubectl create secret generic wordpress-secret \
--from-literal=WORDPRESS_DB_USER=root \
--from-literal=WORDPRESS_DB_PASSWORD=admin \
--from-literal=WORDPRESS_DB_NAME=blog \
--from-literal=WORDPRESS_DB_HOST=dbsvc

# create mysql pod
kubectl apply -f mysql-pod.yaml

# mysql svc
kubectl expose pod mysql --name=dbsvc --port=3306

# create wordpress pod
kubectl apply -f wordpress-pod.yaml

# wordpress nodeport svc
kubectl expose pod wordpress --name=myblog --port=80 --type=NodePort

# Users can visit wordpres
# Check the port of myblog service,
kubectl get svc myblog
# Then user can visit from the browser using VM host ip, 192.168.56.11:port,
192.168.56.12:port
```

## References

---

<https://docs.microsoft.com/en-us/dotnet/architecture/dapr-for-net-developers/the-world-is-distributed>

<https://rubygarage.org/blog/monolith-soa-microservices-serverless>

<https://www.redhat.com/en/topics/containers/kubernetes-architecture>

[https://kubernetes.io/docs/concepts/overview/working-with-objects/\\_print](https://kubernetes.io/docs/concepts/overview/working-with-objects/_print)

<https://kubernetes.io/docs/tasks/access-application-cluster/ingress-minikube/>